

Subqueries & query rewriting

Denis Miginsky



The question of the day

What is the most expensive product for each company?



Subquery

SELECT query expects **tables** under some statements (FROM, JOIN, ...)

SELECT query itself produces **table**-like structure (the application of any operator from the relational algebra is always producing a relation)

Thus, we can use **SELECT**-query instead of the actual table, cannot we?



Subquery example

SQL:

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m,
    (SELECT p.BILL_ID AS BILL
     FROM PRODUCT p, CATEGORY c
     WHERE p.WARE=c.WARE AND c.CLASS='Food') b
WHERE m.BILL_ID=b.BILL;
```

Q: What does this query do?

Q: Can it be rewritten without subqueries?



Pros and cons

Pros:

- Easier to design
- Query is more structured and better reflects the original question
- Some queries could not be written without subqueries

Cons:

- Worse performance in general

Tips:

- Feel free to use subqueries to create the initial query design
- Next, try to rewrite the query with as few subqueries as possible
- Try to keep the subqueries as simple as possible. Move all the complexity to the parent query.



Subqueries under WHERE: IN statement

SQL:

```
SELECT DISTINCT COMPANY
FROM MANUFACTURER
WHERE BILL_ID IN (SELECT BILL_ID
                  FROM PRODUCT
                  WHERE WARE IN (SELECT WARE
                                FROM CATEGORY
                                WHERE CLASS='Food')) ;
```



Limitations for IN statements

IN expects a “list” of values of the same type.

Thus, it expects that the subquery will return the result with the single column, that is compatible with the expression before **IN** by the type.

In general, you can use a tuple on the left and the corresponded projection (by the type) on the right.



Correlated subqueries: EXISTS

SQL:

```
SELECT DISTINCT p.WARE
FROM PRODUCT p
WHERE EXISTS (SELECT 1
              FROM MATERIAL mt1
              WHERE mt1.BILL_ID=p.BILL_ID  --dep. from outer q.
              AND mt1.WARE='Water');
```

Q: What does this query do?

Q: Can it be rewritten without subqueries?



NOT EXISTS

SQL:

```
SELECT DISTINCT p.WARE  
FROM PRODUCT p  
WHERE NOT EXISTS (SELECT 1  
                   FROM MATERIAL mt  
                   WHERE mt.BILL_ID=p.BILL_ID) ;
```

<the same questions here>



Notes on EXISTS

EXISTS checks if the subquery returns at least one row. The structure and values of the row do not matter (so it is the typical pattern to use constant-value projections with value 0 or 1).

Usually it can be rewritten with **JOIN** and **DISTINCT** in the projection.

NOT EXISTS is opposite to **EXISTS**.

Usually it can be rewritten with **OUTER JOIN** and the explicit **IS NULL** check, or with **EXCEPT**.



Lexical scope in subqueries

The subquery forms its own lexical scope that is nested to the scope of the parent query:

- The parent query can see only the result of the subquery, but not its local “variables”.
- The **renaming** (**AS**) operation is essential for subqueries (unlike plain queries where the renaming is purely cosmetic)
- Correlated subqueries can appear in selection and projection sections (under **SELECT**, **WHERE**, **HAVING**, but not under **FROM**)
- The correlated subquery can see everything from the parent query that is available for the statement where it appears (consider **HAVING**/**GROUP BY** limitations for instance)



Question

For each product we need all the materials. There must be one row per product ware.

Looks familiar?

Can we rewrite this without GROUP BY?

What if we need the list of materials to be ordered?



Correct query

SQL:

```
SELECT p.WARE,  
       (SELECT GROUP_CONCAT(WARE)  
        FROM (SELECT DISTINCT mt1.WARE  
              FROM MATERIAL mt1, PRODUCT p1  
              WHERE mt1.BILL_ID=p1.BILL_ID  
                   AND p1.WARE=p.WARE  
              ORDER BY mt1.WARE) )  
FROM (SELECT DISTINCT WARE FROM PRODUCT) p  
ORDER BY p.WARE;
```



Is it an aggregation?

```
SELECT WARE, PRICE  
FROM PRODUCT  
ORDER BY PRICE DESC  
LIMIT 1;
```

Q: What is the meaning of this query?

Q: Can we do the same for each company?



Answer to the question of the day

```
SELECT DISTINCT m.COMPANY,  
  (SELECT p1.WARE  
   FROM PRODUCT p1, MANUFACTURER m1  
   WHERE p1.BILL_ID=m1.BILL_ID  
        AND m1.COMPANY=m.COMPANY  
   ORDER BY p1.PRICE DESC  
   LIMIT 1) AS EXP_WARE,  
  (SELECT p1.PRICE  
   FROM PRODUCT p1, MANUFACTURER m1  
   WHERE p1.BILL_ID=m1.BILL_ID  
        AND m1.COMPANY=m.COMPANY  
   ORDER BY p1.PRICE DESC  
   LIMIT 1) AS PRICE  
FROM MANUFACTURER m;
```

Problem: a little bit bulky, there are two almost identical queries

Problem: the performance is not perfect



More efficient variants

```
SELECT DISTINCT m.COMPANY,  
               ... AS EXP_WARE,  
               ... AS PRICE  
FROM (SELECT DISTINCT COMPANY FROM MANUFACTURER) m;
```

```
SELECT m.COMPANY,  
       ... AS EXP_WARE,  
       ... AS PRICE  
FROM (SELECT DISTINCT COMPANY FROM MANUFACTURER) m;
```

```
SELECT m.COMPANY,  
       ... AS EXP_WARE,  
       ... AS PRICE  
FROM MANUFACTURER m  
GROUP BY m.COMPANY;
```



Subqueries as values

```
SELECT DISTINCT m.COMPANY  
FROM MANUFACTURER m  
WHERE (SELECT COUNT(BILL_ID) FROM MANUFACTURER m1  
       WHERE m1.COMPANY=m.COMPANY) > 2;
```

Problem: the performance!



Other variants

```
SELECT m.COMPANY
FROM (SELECT DISTINCT COMPANY FROM MANUFACTURER) m
WHERE (SELECT COUNT(BILL_ID) FROM MANUFACTURER m1
      WHERE m1.COMPANY=m.COMPANY) > 2;
```

```
SELECT m.COMPANY,
      (SELECT COUNT(BILL_ID) FROM MANUFACTURER m1
      WHERE m1.COMPANY=m.COMPANY) AS CNT
FROM (SELECT DISTINCT COMPANY FROM MANUFACTURER) m
WHERE CNT > 2;
```

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m
GROUP BY m.COMPANY
HAVING COUNT(BILL_ID)>2; --the fastest one
```



COUNT vs EXISTS: performance

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m
WHERE EXISTS (SELECT 1 FROM MANUFACTURER m1, MATERIAL mt1
              WHERE mt1.BILL_ID=m1.BILL_ID
              AND m.COMPANY=m1.COMPANY);
```

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m
WHERE (SELECT COUNT() FROM MANUFACTURER m1, MATERIAL mt1
      WHERE mt1.BILL_ID=m1.BILL_ID
      AND m.COMPANY=m1.COMPANY)>0;
```

Q: Which one is better?



The best variant: no COUNT/EXISTS

```
SELECT DISTINCT m.COMPANY  
FROM MANUFACTURER m  
JOIN MATERIAL mt  
ON mt.BILL_ID=m.BILL_ID;
```



COUNT vs NOT EXISTS: performance

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m
WHERE NOT EXISTS (SELECT 1 FROM MANUFACTURER m1, MATERIAL mt1
                  WHERE mt1.BILL_ID=m1.BILL_ID
                  AND m.COMPANY=m1.COMPANY);
```

```
SELECT DISTINCT m.COMPANY
FROM MANUFACTURER m
WHERE (SELECT COUNT() FROM MANUFACTURER m1, MATERIAL mt1
      WHERE mt1.BILL_ID=m1.BILL_ID
      AND m.COMPANY=m1.COMPANY)=0;
```

Q: Which one is better?



The best variant

```
SELECT DISTINCT m.COMPANY  
FROM MANUFACTURER m  
LEFT JOIN MATERIAL mt  
ON mt.BILL_ID=m.BILL_ID  
WHERE mt.BILL_ID IS NULL;
```

Q: Is it correct? Is it equivalent to the previous ones?



Pros and cons

| | Performance | Flexibility |
|--------|-------------|-------------|
| COUNT | The worst | The best |
| EXISTS | Bad | Good |
| JOIN | The best | The Worst |



Notes on performance and flexibility

- Being closer to the formal model (relational algebra) means better performance: **plain query** vs **nested queries**
- Better flexibility means worse performance: **EXISTS/COUNT** vs **JOIN**
- Highly specialized/higher level operations provide more information to the query planner, and it can make a better plan: **EXISTS** vs **COUNT**, **equijoin** vs **θ -join**



Typical situations where subqueries are essential

- The aggregation should be used, and more complex manipulations with groups are required than standard aggregators can provide:
 - Ordering
 - Sampling
 - ...
- Set operations should be used with subsequent manipulations with results:
 - Aggregation
 - Join
 - ...
- The appropriate level of query flexibility is necessary: one needs to change conditions in a parent query and subqueries relatively independently.
- ...

