

# Query planning

**Denis Miginsky**



# The question of the day

For the following query

```
SELECT DISTINCT p.WARE, m.COMPANY
FROM MANUFACTURER m, PRODUCT p, CATEGORY c
WHERE c.CLASS='Raw food' AND m.BILL_ID=p.BILL_ID
      AND c.WARE=p.WARE
ORDER BY p.WARE ASC
LIMIT 10
```

what is the asymptotic complexity?



# Possible query plan (1)

Let's consider SQL as functional language

**FILTER**, **MAP**, **TAKE**, etc. are as in any functional language

$X \text{ CROSS\_PROD } Y = [(x,y) \mid x \leftarrow X, y \leftarrow Y]$  -- Haskell-like

## Plan (pseudo-code)

MANUFACTURER **CROSS\_PROD** PRODUCT

-> **CROSS\_PROD** CATEGORY

-> **FILTER** c.CLASS='Raw food'

-> **FILTER** m.BILL\_ID=p.BILL\_ID

-> **FILTER** c.WARE=p.WARE

-> **SORT\_BY** p.WARE

-> **DISTINCT**

-> **MAP** (p.WARE, m.COMPANY)

-> **TAKE** 10

## Let:

M=size(MANUFACTURER) ~1000

P=size(PRODUCT) ~2000

C=size(CATEGORY) ~20

**Q:** What is the asymptotic complexity of this plan?



# Complexity

Plan (pseudo-code)	Complexity	Cardinality
MANUFACTURER <b>CROSS_PROD</b> PRODUCT	$O(M \cdot P)$	$M \cdot P$
-> <b>CROSS_PROD</b> CATEGORY	$O(M \cdot P \cdot C)$	$M \cdot P \cdot C$
-> <b>FILTER</b> c.CLASS='Raw food'	$O(M \cdot P \cdot C)$	$M \cdot P \cdot C' \mid C' < C$
-> <b>FILTER</b> m.BILL_ID=p.BILL_ID	$O(M \cdot P \cdot C')$	$P \cdot C'$
-> <b>FILTER</b> c.WARE=p.WARE	$O(P \cdot C')$	$\sim P \cdot C' / C < P$
-> <b>SORT_BY</b> p.WARE	$O(P \cdot C' / C \cdot \log(P \cdot C' / C))$	$\sim P \cdot C' / C$
-> <b>MAP</b> (p.WARE, m.COMPANY)	$O(P \cdot C' / C)$	$\sim P \cdot C' / C$
-> <b>DISTINCT</b>	$O((P \cdot C' / C)^2)$	$\sim P \cdot C' / C$
-> <b>TAKE</b> 10	10	10
<b>TOTAL:</b>	<b><math>O(M \cdot P \cdot C)</math></b>	<b>10</b>

What options are available to improve the performance?



# Option 1: algebraic properties

Fortunately, SQL is **not** a functional language. In fact, it is a level higher than a functional language. This means that:

1. Query can be compiled to the functional program in many ways
2. Since SQL is an implementation of relational algebra, more properties could be in use besides just a “crude”  $\beta$ -reduction
3. Such properties are the algebraic properties of the relational algebra’s operations: associativity and commutativity.



# Some properties of relational algebra

Not so formally speaking:

- Joins are associative:  $\mathbf{A} \bowtie \mathbf{B}$  and  $\mathbf{A} \bowtie \mathbf{C}$  are associative (assuming any type of the inner or even the outer join)
- Inner joins are commutative (in fact, because formally they are not because we have to “revert” the  $\theta$ -join predicate)
- Selections are associative with each other and even with joins (until there are enough attributes on the particular stage to apply the filter)

**Thus:** there are multiple variants to reorder the plan’s elements.



# Goals and hints for optimization

- **Goal:** achieve the total complexity as close as possible to the final cardinality of the query
- **Sub-goal:** keep the complexity as small as possible of each stage
- **Rule:** keep the cardinality as small as possible
- **Hint:** place the filters as early as possible (they always reduce the cardinality)
- **Hint:** place the joins (and other operations) with the lower cardinality as early as possible



# Better plan (2)

Plan (pseudo-code)	Complexity	Cardinality
CATEGORY <b>FILTER</b> c.CLASS='Raw food'	$O(C)$	$C' < C$
-> <b>CROSS_PROD</b> PRODUCT	$O(P * C')$	$P * C'$
-> <b>FILTER</b> c.WARE=p.WARE	$O(P * C')$	$\sim P * C' / C < P$
-> <b>CROSS_PROD</b> MANUFACTURER	$O(M * P * C' / C)$	$\sim M * P * C' / C$
-> <b>FILTER</b> m.BILL_ID=p.BILL_ID	$O(M * P * C' / C)$	$\sim P * C' / C$
-> <b>SORT_BY</b> p.WARE	$O(P * C' / C * \log(P * C' / C))$	$\sim P * C' / C$
-> <b>MAP</b> (p.WARE, m.COMPANY)	$O(P * C' / C)$	$\sim P * C' / C$
-> <b>DISTINCT</b>	$O((P * C' / C)^2)$	$\sim P * C' / C$
-> <b>TAKE</b> 10	10	10
<b>TOTAL:</b>	<b><math>O(M * P * C' / C)</math></b>	<b>10</b>

Much better complexity! However, it is terrible yet (any non-linear is terrible).





# Option 2: better algorithms

- There are multiple algorithms for the special cases of join (Cartesian product with the filter is far from the best)
- There are better algorithms for the special cases of selection rather than the crude **filter**
- Better algorithms require specific properties of the data organization



# Indices

In general, the index is the special data structure that contains all the entries from the original table (but, probably, not the whole entries), that helps with searching.

**Tree index** – the most common index, implementing multimap **index\_attr** → **row\_id** (by a sort of balanced tree, usually a variant of B-tree) where **index\_attr** – the attribute of the choice, **row\_id** – internal DB identifier of the row (with the lookup complexity of **O(1)**).

The tree index has the following properties:

- The lookup and the insertion costs are **O(log(N))** per element
- The iteration cost is **O(1)** per element
- All the entries are ordered by **index\_attr**



# Join algorithms

- Nested loop join
- Hash join
- Merge join



# Nested loops join

Nested loops join of **TAB1** and **TAB2** on predicate **P**:

For each **row1** from **TAB1**, for each **row2** from **TAB2**:  
    **result** += (**row1**, **row2**) when **P(row1, row2)**

Assuming **M=size(TAB1)**, **N=size(TAB2)**

## Pros:

- Any predicate is supported
- No prerequisites on the data organization
- Preserves order of **TAB1**

## Cons:

- Complexity  **$O(M*N)$**



# Hash join in DB

Hash join of **TAB1** and **TAB2** on attributes **a1(TAB1)** and **a2 (TAB2)**:

Assuming **IDX2=index(TAB2, a2)** (hash-table originally)

For each **row1** from **TAB1** (called the leading table):  
    **result** += **IDX2[a1(row1)]**

## Pros:

- Complexity  **$O(N \cdot \log(N)) + O(M \cdot \log(N))$**  (no first part when the index is pre-built)
- Preserves the order of **TAB1**

## Cons:

- Limited predicates are supported (= for traditional hash-join, also <, > for DB-version)



# Original vs DB hash-joins

Hash-join is usually recommended to be used when at least one table is original and has pre-built index.

When both tables have indices, a smaller one is preferred as the leading one.

Complexity is asymmetrical for TAB1 and TAB2.

	Original	DB
<b>Index type</b>	hash-table	tree-map
<b>Recommended leading table</b>	largest	smallest
<b>Predicates</b>	=	=, >, <



# Merge-join

Merge join of **TAB1** and **TAB2** on attributes **a1(TAB1)** and **a2 (TAB2)**:

Sort **TAB1** and **TAB2**.

Next use the algorithm similar to merging ordered arrays in merge-sort.

## Pros:

- Complexity  $O(M \cdot \log(M)) + O(N \cdot \log(N)) + O(M) + O(N)$  (no first part when TAB1 and TAB2 either pre-sorted or have indices)
- The result is sorted

## Cons:

- Limited predicates are supported ( $=$ ,  $<$ ,  $>$ )



# Plan 2

Plan (pseudo-code)	Complexity	Cardinality
CATEGORY <b>FILTER</b> c.CLASS='Raw food'	$O(C)$	$C' < C$
-> <b>NL_JOIN</b> PRODUCT <b>ON</b> c.WARE=p.WARE	$O(P * C)$	$\sim P * C' / C < P$
-> <b>NL_JOIN</b> MANUFACTURER <b>ON</b> m.BILL_ID=p.BILL_ID	$O(M * P * C' / C)$	$\sim P * C' / C$
-> <b>SORT_BY</b> p.WARE	$O(P * C' / C * \log(P * C' / C))$	$\sim P * C' / C$
-> <b>MAP</b> (p.WARE, m.COMPANY)	$O(P * C' / C)$	$\sim P * C' / C$
-> <b>DISTINCT</b>	$O((P * C' / C)^2)$	$\sim P * C' / C$
-> <b>TAKE</b> 10	10	10
<b>TOTAL:</b>	$O(M * P * C' / C)$	<b>10</b>

The same plan, just re-written with nested loops join





# Plan 3

Plan (pseudo-code)	Complexity	Cardinality
CATEGORY <b>FILTER</b> c.CLASS='Raw food'	$O(C)$	$C' < C$
-> <b>HASH_JOIN</b> PRODUCT <b>INDEX BY</b> WARE ON c.WARE=p.WARE	$O(P \cdot C' / C)$	$\sim P \cdot C' / C < P$
-> <b>HASH_JOIN</b> MANUFACTURER <b>INDEX BY</b> BILL_ID ON m.BILL_ID=p.BILL_ID	$O(P \cdot C' / C)$	$\sim P \cdot C' / C$
-> <b>SORT_BY</b> p.WARE	$O(P \cdot C' / C \cdot \log(P \cdot C' / C))$	$\sim P \cdot C' / C$
-> <b>MAP</b> (p.WARE, m.COMPANY)	$O(P \cdot C' / C)$	$\sim P \cdot C' / C$
-> <b>DISTINCT</b>	$O((P \cdot C' / C)^2)$	$\sim P \cdot C' / C$
-> <b>TAKE</b> 10	10	10
<b>TOTAL:</b>	<b><math>O((P \cdot C' / C)^2)</math></b>	<b>10</b>

Same as plan 2, but with a better join algorithm



# Option 3: laziness

Why shall we process the entire result if we need only a few top rows?

Which of the following operations could be lazy?

- Nested loops join
- Hash join
- Merge join
- Selection/filter
- Projection
- Ordering
- Distinct
- Count



# Plan 3 with laziness

Plan (pseudo-code)	Complexity	Cardinality
CATEGORY <b>FILTER</b> c.CLASS='Raw food'	$O(C)$	$C' < C$
-> <b>HASH_JOIN</b> PRODUCT <b>INDEX BY</b> WARE ON c.WARE=p.WARE	$O(P \cdot C' / C)$	$\sim P \cdot C' / C < P$
-> <b>HASH_JOIN</b> MANUFACTURER <b>INDEX BY</b> BILL_ID ON m.BILL_ID=p.BILL_ID	$O(P \cdot C' / C)$	$\sim P \cdot C' / C$
-> <b>SORT_BY</b> p.WARE	$O(P \cdot C' / C \cdot \log(P \cdot C' / C))$	$\sim P \cdot C' / C$
-> <b>MAP</b> (p.WARE, m.COMPANY)	$\sim 10$	$\sim 10$
-> <b>DISTINCT</b>	$\sim 50$	10
-> <b>TAKE</b> 10	10	10
<b>TOTAL:</b>	$O(P \cdot C' / C \cdot \log(P \cdot C' / C))$	<b>10</b>

The complexity is better, non-linear still.

The sorting instruction breaks the laziness. Can we fix this?



# Plan 4

Plan (pseudo-code)	Complexity	Cardinality
CATEGORY INDEX BY WARE FILTER c.CLASS='Raw food'	~20	$C' < C$
-> MERGE_JOIN PRODUCT INDEX BY WARE ON c.WARE=p.WARE	$\sim 100 = \sim 10 * C / C'$	~10
-> HASH_JOIN MANUFACTURER INDEX BY BILL_ID ON m.BILL_ID=p.BILL_ID	~10	~10
-> MAP (p.WARE, m.COMPANY)	~10	~10
-> DISTINCT	~50	10
-> TAKE 10	10	10
<b>TOTAL:</b>	<b>~200</b>	<b>10</b>

That is the best plan I have for now. Can you do it better?

